

Optimal Non-Adaptive Group Testing under Markovian Infection Models

Aditya Narayan Ravi and Ilan Shomorony
 Department of Electrical and Computer Engineering
 University of Illinois Urbana-Champaign, Illinois, USA
 Email: {anravi2, ilans}@illinois.edu

Abstract—THIS PAPER IS ELIGIBLE FOR THE STUDENT PAPER AWARD. We consider group testing under a Markov infection model, in the sparse-infection regime where the expected number of infected items scales as $O(n^\theta)$ for some $\theta \in (0, 1)$. In this setting, we exactly characterize the optimal number of non-adaptive tests. The resulting capacity generalizes the independent-prior capacity results of Coja-Oghlan et al. [1] to a nontrivial setting with correlated priors. Our results imply that Markovian correlations reduce the non-adaptive testing capacity in sufficiently sparse regimes. However, they also identify a threshold in the sparsity parameter θ above which correlations provide no improvement in the non-adaptive testing capacity.

I. INTRODUCTION

We consider a *correlated* group testing problem in which the infection vector is generated by a Markov source. We investigate the non-adaptive [2] testing capacity of this model.

Group testing was introduced in [3] as a strategy for reducing the number of tests needed to estimate disease prevalence in large populations, when infections are sparse. A classical formulation is the independent combinatorial model, in which a population of size n contains exactly k infected individuals, chosen uniformly at random from all subsets of size k . Formally, the infection vector is denoted by $U^n \in \{0, 1\}^n$, where $U_i = 1$ indicates that item i is infected. A non-adaptive test design is represented by a binary matrix $X \in \{0, 1\}^{T \times n}$, where all tests are fixed before observing any outcomes, and $X_{ti} = 1$ indicates that item i is included in test t . The corresponding noiseless test outcomes are given by

$$Y_t = \mathbf{1} \left\{ \sum_{i=1}^n X_{ti} U_i > 0 \right\}, \quad t = 1, \dots, T.$$

Thus, a test returns positive if it contains at least one infected item. The goal is to construct an estimator $\hat{U}^n = \hat{U}^n(X, Y)$ that recovers U^n from the test design and outcome vector.

This naturally led to interest in characterizing the fundamental limits on the number of tests required for reliable recovery of the infected individuals. The capacity of group testing under such independent combinatorial infection models has since been studied extensively [2], [4]–[8]. In particular, the non-adaptive capacity under the independent combinatorial model was resolved by Coja-Oghlan et al. in [1], [9].

A natural motivation for the correlated setting is that independence rarely reflects real-world infection processes. Infection vectors are often better modeled by correlated sources, where an item’s infection status depends on the infection status

of other items, rather than being independent across items. For example, individuals in the same household, neighborhood, or community are more likely to be infected together, producing clustered or bursty infection patterns [10]–[12]. Similar correlation structures also arise in biological testing applications, such as spatial transcriptomics [13], where co-regulated genes may be activated together. These settings motivate a theory of non-adaptive group testing under correlated infection priors.

Many recent works have sought to relax the independence assumption. One line of work studies community-based infection models, in which infections spread over predefined graph structures [14]–[18]. Other works consider alternative correlated priors, including Markovian infection models [19], [20], non-uniform priors [21] and hypergraph models [22].

The optimal number of tests in the non-adaptive setting remains poorly understood for most correlated models. For example, [14], [15] characterize rates under a constant, non-vanishing false-positive rate, but do not provide a sharp threshold. Similarly, [19] focuses on computationally efficient decoders whose testing rates remain strictly suboptimal.

As a concrete instance of a correlated setting with this technical obstacle, we consider the Markovian infection model introduced in [19], considered in Figure 1. In this model, the infection vector U^n is generated by a two-state Markov chain, initialized according to its stationary distribution. The transition probabilities are α_n and β , where $\alpha_n = k'n^{\theta-1}$ for some constant k' and $\beta \in (0, 1)$ is a constant. Thus, infections occur in bursts: once the chain enters state 1, it remains infected for a geometrically distributed number of consecutive positions with mean $1/\beta$. The stationary marginal infection probability scales as $q_n \asymp n^{\theta-1}$, and hence the expected number of infected items scales as $nq_n \asymp n^\theta$, consistent with the sparse regime we consider in this paper.




Fig. 1. The 2-state Markov chain that generates the infection vector U^n




Fig. 2. Normalized number of tests required for exact recovery under the independent prior and correlated Markov priors with different values of β . Correlations reduce the testing threshold in sparse regimes, while all curves coincide with the independent threshold $\theta/\ln 2$ once θ is sufficiently large.

Under this model, we establish a sharp testing threshold T_{cap} , where

$$T_{\text{cap}} = \max \left\{ \beta(1 - \theta), \frac{\theta}{\ln 2} \right\} nq \log n,$$

where $\log$ is base 2. Reliable recovery of U^n is asymptotically possible if T is above this threshold and impossible below it.

Figure 2 compares the normalized number of tests required under the independent and Markovian infection models. In particular, when $\theta > \ln 2 / (1 + \ln 2)$, the Markovian correlation does not improve the testing capacity relative to the independent case. As $\beta \rightarrow 1$, the Markovian model approaches the independent-prior setting, and the resulting threshold recovers the corresponding independent group testing threshold.

The main technical contribution is a new source-dependent typical set, which captures the high-probability infection vectors generated by the prior. Unlike in the classical combinatorial model, competitors at a fixed overlap cannot be counted by binomial arguments. We develop sharp large-deviation and Sanov-type counting bounds for these typical sequences and combine them with test-design collision probabilities. This yields a decoupling principle: the error probability separates into a source-complexity term, controlling the number of plausible competitors, and a test-design term, controlling the probability that such competitors match the observed outcomes. This decoupling enables the exact characterization of the optimal non-adaptive testing threshold.

II. PROBLEM FORMULATION AND PRELIMINARIES

We consider the Markovian correlated group testing problem studied in [19], in which n distinct items, indexed by $\{1, 2, \dots, n\}$, are assigned binary labels $U_i \in \{0, 1\}$.

The two-state Markov chain has a state space $S = \{0, 1\}$ and transition probabilities $p_{0,1} = \alpha$ and $p_{1,0} = \beta$, as shown in Figure 1. The infection vector U^n is generated by initializing the chain in its stationary distribution

$$(q, 1 - q) := \left(\frac{\alpha}{\alpha + \beta}, \frac{\beta}{\alpha + \beta} \right), \quad (1)$$

where $q := q_n$. Then we record n consecutive states.

Let $T := T(n)$ denote the number of tests. A non-adaptive test design is specified by a binary matrix $X \in \{0, 1\}^{T \times n}$, chosen before any test outcomes are observed, where $X_{t,i} = 1$ indicates that item i is included in test t . The noiseless outcome vector $Y = (Y_1, \dots, Y_T)$ is then given by

$$Y_t = \mathbf{1} \left\{ \sum_{i=1}^n X_{t,i} U_i > 0 \right\}, \quad t = 1, \dots, T. \quad (2)$$

Equivalently, $Y_t = 1$ if t contains at least one infected item. A decoding rule G then maps (X, Y) to an estimate $\hat{U}^n \in \{0, 1\}^n$. Let $\mathcal{E} = \{U^n \neq \hat{U}^n\}$ denote the error event.

The objective is to construct a sequence of non-adaptive test matrices and associated decoders such that $\Pr(\mathcal{E}) \rightarrow 0$ as $n \rightarrow \infty$. Following [4], we study the regime in which the number of tests scales on the order of $nq_n \log n$. Based on this scaling, we define an achievable number of normalized tests as follows:

Definition 1 (Achievability): A normalized number of tests τ is said to be *achievable* if there exists a sequence of non-adaptive test matrices $X_n \in \{0, 1\}^{T(n) \times n}$ and decoding rules $g_n : \{0, 1\}^{T(n) \times n} \times \{0, 1\}^{T(n)} \rightarrow \{0, 1\}^n$ such that, with $\hat{U}^n = g_n(X_n, Y)$,

$$\lim_{n \rightarrow \infty} \frac{T(n)}{nq_n \log n} = \tau, \quad (3)$$

and $\Pr(\mathcal{E}) \rightarrow 0$ as $n \rightarrow \infty$. The optimal normalized number of tests is defined as

$$\tau^* := \inf \{ \tau : \tau \text{ is achievable} \}. \quad (4)$$

The infimum is taken over all sequences of non-adaptive test designs $\{X^{(n)}\}_{n \geq 1}$ and all corresponding decoders $\{g_n\}_{n \geq 1}$. Equivalently, for any $\tau < \tau^*$, there is no sequence of test designs and decoders for which the error probability vanishes as $n \rightarrow \infty$. We note that while this definition is stated differently from the capacity formulation in [4], the two definitions characterize the same fundamental threshold.

Following the group testing literature, we work in the sparse regime [4], where the expected number of infected items satisfies $nq_n = \Theta(n^\theta)$, $\theta \in (0, 1)$. We specify

$$\alpha = \frac{k'n^\theta}{n}, \quad (5)$$

with $k' > 0$ constant and $\beta \in (0, 1)$ fixed. Then

$$q_n = \frac{\alpha}{\alpha + \beta} = \frac{k'}{\beta} n^{\theta-1} + o(n^{\theta-1}). \quad (6)$$

Writing $k := k'/\beta$, we use the shorthand $q_n = kn^\theta/n$, omitting lower order terms, since they do not affect results.

We introduce a notion of typical infection vectors tailored to the sparse Markov source. This construction is based on a re-normalized entropy rate, which captures the source complexity.

Lemma 1 (Sparse entropy rate): For the Markov infection source described above,

$$\frac{H(U^n)}{nq_n \log n} = \beta(1 - \theta) + o(1).$$

We then define a typical set in terms of this entropy rate.

Definition 2 (ϵ -typical set): For $\epsilon > 0$, define

$$T_\epsilon^{(n)} := \left\{ u^n : \left| -\frac{1}{nq_n \log n} \log P_{U^n}(u^n) - \beta(1 - \theta) \right| \leq \epsilon, \right. \\ \left. \left| \frac{1}{n} \sum_{i=1}^n u_i - q_n \right| \leq \epsilon \right\}. \quad (7)$$

The first condition states that every sequence in $T_\epsilon^{(n)}$ has probability close to the typical probability scale of the source. The second condition enforces the sparse weight constraint: every typical sequence has approximately nq_n infected entries.

Lemma 2 (Typical-set concentration): Consider the Markov infection source described above, and let $T_\epsilon^{(n)}$ be the ϵ -typical set in Definition 2. Then, for every fixed $\epsilon > 0$,

$$\mathbb{P}(U^n \in T_\epsilon^{(n)}) \rightarrow 1 \quad \text{as } n \rightarrow \infty.$$

The proofs of Lemmas 1 and 2 are provided in a longer version of this paper.

Notation. Throughout, $\log$ denotes base-2 logarithm and $\ln$ denotes the natural logarithm. We use the standard asymptotic notation $o(\cdot)$ and $\Theta(\cdot)$ as $n \rightarrow \infty$. Throughout, for any positive integer m , we use the notation $[m] := \{1, 2, \dots, m\}$.

III. TEST DESIGN AND MAIN RESULTS

Before stating the full result, we describe the test matrix design and decoding rule used throughout the achievability analysis. The construction largely follows the near-constant-column design and consistency decoder considered in [9].

Recall that $X \in \{0, 1\}^{T \times n}$ denotes the non-adaptive test matrix. Fix a design parameter $\nu > 0$ and set

$$\Delta := \frac{\nu T}{nq_n}.$$

For each item $i \in [n]$, independently choose Δ test indices uniformly at random from $[T]$, with replacement, and set $X_{t,i} = 1$ if item i is assigned to test t at least once.

For any $u^n \in \{0, 1\}^n$, let $Y(u^n)$ denote the noiseless outcome vector induced by u^n under X . Given X and Y , the decoder searches for vectors in the typical set that are consistent with the observed outcomes:

$$\mathcal{C}(Y) := \left\{ u^n \in T_\epsilon^{(n)} : Y(u^n) = Y \right\}.$$

If $\mathcal{C}(Y) = \{u^n\}$ is a singleton, the decoder returns $\hat{U}^n = u^n$. Otherwise the decoder declares an error.

We now state our main theorem which characterizes the optimal normalized number of tests.

Theorem 3 (Optimal normalized number of tests): Consider the sparse regime $nq_n = n^\theta$, where $\theta \in (0, 1)$, and let τ^* be the optimal normalized number of tests. Then

$$\tau^* = \max \left\{ \beta(1 - \theta), \frac{\theta}{\ln 2} \right\}.$$

The threshold has two distinct contributions. The term $\beta(1 - \theta)$ is the source-complexity term, which reflects the number of tests needed to distinguish the true infection vector from the collection of other typical infection vectors. The term $\theta / \ln 2$ is a high-overlap indistinguishability term, which

arises from competitors that differ from the true vector in only a small number of infected coordinates and are therefore difficult to separate through pooled tests. A related obstruction was identified by [9] in the independent prior setting. The key technical challenge for this term is to show that the same obstruction persists under the Markov source, despite the correlations between infection coordinates and the non-uniform structure of the source. We prove the upper bound in Section IV and the lower bound in Section V.

IV. ACHIEVABILITY

We now outline the achievability argument under the constant test design and the typicality-constrained consistency decoder. Define the overall error event $\mathcal{E} := \mathcal{E}_0 \cup \mathcal{E}_1$, where

$$\mathcal{E}_0 = \left\{ U^n \notin T_\epsilon^{(n)} \right\}, \\ \mathcal{E}_1 = \left\{ \exists v^n \in T_\epsilon^{(n)}, v^n \neq U^n : Y(v^n) = Y(U^n) \right\}. \quad (8)$$

By the axiom of probability,

$$\mathbb{P}(\mathcal{E}) = \mathbb{P}(\mathcal{E}_0) + \mathbb{P}(\mathcal{E}_1 | \mathcal{E}_0^c). \quad (9)$$

Since $T_\epsilon^{(n)}$ is a typical set satisfying Lemma 2, for all sufficiently large n , we have

$$\mathbb{P}(\mathcal{E}) \leq \mathbb{P}(\mathcal{E}_1 | U^n \in T_\epsilon^{(n)}) + \epsilon. \quad (10)$$

We now bound this by grouping competitors v^n according to their “infected-overlap” with the realized sequence U^n . For two fixed sequences $u^n, v^n \in \{0, 1\}^n$, let

$$l(u^n, v^n) := \sum_{i=1}^n \mathbf{1}\{u_i = 1, v_i = 1\} \quad (11)$$

denote the number of coordinates at which both sequences are equal to 1. Recall $k = nq$, and define

$$\lambda(u^n, v^n) := \frac{l(u^n, v^n)}{k + \epsilon}, \quad (12)$$

where λ is the normalized overlap. For each value of λ , let

$$N(\lambda | u^n) := \left| \left\{ v^n \in T_\epsilon^{(n)} \setminus \{u^n\} : \frac{l(u^n, v^n)}{k + \epsilon} = \lambda \right\} \right|,$$

denote the number of sequences $v^n \in T_\epsilon^{(n)}$ with $v^n \neq u^n$ and normalized overlap λ with u^n . Let

$$\lambda_n^* := 1 - \frac{1}{\log n}$$

be the overlap threshold used to separate the small-overlap and large-overlap regimes. We now decompose the error event according to the normalized overlap between the true sequence u^n and a competing sequence v^n . Let

$$p_T(u^n) := \mathbb{P}(U^n = u^n | U^n \in T_\epsilon^{(n)}).$$

For overlap levels $\lambda \leq \lambda_n^*$, we union bound over all competitors in that overlap class. The number of such competitors is denoted by $N(\lambda | u^n)$. All competitors with overlap larger than λ_n^* are collected into a single large-overlap event. Thus,

$$\mathbb{P}(\mathcal{E}_1 | U^n \in T_\epsilon^{(n)}) = \sum_{u^n \in T_\epsilon^{(n)}} p_T(u^n) \mathbb{P}(\mathcal{E}_1 | U^n = u^n)$$

$$\leq \sum_{u^n \in T_\epsilon^{(n)}} p_T(u^n) \left[\sum_{\lambda \leq \lambda_n^*} N(\lambda | u^n) \mathbb{P}(\mathcal{C}(u^n, v^n) | \mathcal{O}_\lambda(u^n, v^n)) + \mathbb{P}(\mathcal{C}_{>\lambda_n^*}(u^n) | U^n = u^n) \right]. \quad (13)$$

Here, in the first term, v^n denotes an arbitrary typical competitor in the overlap class

$$\mathcal{O}_\lambda(u^n, v^n) := \{u^n, v^n \in T_\epsilon^{(n)} : \lambda(u^n, v^n) = \lambda\},$$

and

$$\mathcal{C}(u^n, v^n) := \{Y(v^n) = Y(u^n)\}$$

is the event that this fixed competitor produces the same test outcome as u^n . The large-overlap event is

$$\mathcal{C}_{>\lambda_n^*}(u^n) := \left\{ \exists v^n \in T_\epsilon^{(n)} \setminus \{u^n\} : \begin{array}{l} \lambda(u^n, v^n) > \lambda_n^*, \\ Y(v^n) = Y(u^n) \end{array} \right\}.$$

We control the first term in (13) by combining two bounds. First, we bound $N(\lambda | u^n)$, the number of typical competitors with overlap λ . Second, we bound the probability that one fixed competitor in this overlap class produces the same test outcome as u^n . We control the second term directly by showing that the near-constant-column design rules out all typical competitors with overlap larger than λ_n^* . We now state the corresponding bounds and give proof sketches.

Lemma 3 ($N(\lambda)$ bound): Fix $\epsilon > 0$ and let $u^n \in T_\epsilon^{(n)}$. Then for any fixed $u^n \in T_\epsilon^{(n)}$,

$$N(\lambda | u^n) \leq 2^{nq \log n [\beta(1-\theta)(1-\lambda) + \epsilon + o_n(1)]}.$$

Proof Sketch: Fix $u^n \in T_\epsilon^{(n)}$ and condition on an overlap level λ . A competing sequence v^n with $\lambda(u^n, v^n) = \lambda$ must agree with u^n on roughly λnq infected coordinates and introduce roughly $(1-\lambda)nq$ new infected coordinates outside the support of u^n .

Under the Markov source, infected coordinates occur in runs. The number of infected runs in a typical sequence is approximately βnq , because each infected run ends through a $1 \rightarrow 0$ transition, which occurs with probability β . Thus, replacing a $(1-\lambda)nq$ new infections corresponds, up to lower-order effects, to replacing $\beta(1-\lambda)nq$ infected runs.

Each new infected run must begin with a $0 \rightarrow 1$ transition. Since such a transition occurs with probability $\alpha_n = \Theta(n^{\theta-1})$, the number of admissible places in the competitor for the start of new infected runs is $\approx \frac{1}{\alpha_n} = \Theta(n^{1-\theta})$.

Since there are at most $\beta(1-\lambda)nq$ new infected runs, the number of typical competitors with overlap λ is bounded by

$$N(\lambda | u^n) \leq (n^{1-\theta})^{\beta(1-\lambda)nq}$$

up to subexponential factors. Equivalently,

$$N(\lambda | u^n) \leq 2^{nq \log n [\beta(1-\theta)(1-\lambda) + \epsilon + o_n(1)]}.$$

The full proof is available in the longer version of this paper.

Theorem 4 (Small-overlap probability): Fix $u^n, v^n \in T_\epsilon^{(n)}$ with $\lambda(u^n, v^n) = \lambda \leq \lambda_n^*$. Under the above test design,

$$\mathbb{P}(Y(v^n) = Y(u^n) | \mathcal{O}_\lambda(u^n, v^n)) \leq 2^{-T(1-\lambda)[1-\epsilon-o_n(1)]}.$$

Proof Sketch: Here we give a sketch for an i.i.d. Bernoulli test

design, where each item is included in each test independently with probability $p = \nu/k$. The near-constant-column design creates mild dependencies between rows. These dependencies are handled in the full proof and do not change the final result.

Fix $u^n, v^n \in T_\epsilon^{(n)}$ with normalized overlap λ . For clarity, ignore the vanishing typicality deviations and assume

$$|\text{supp}(u^n)| = |\text{supp}(v^n)| = k, |\text{supp}(u^n) \cap \text{supp}(v^n)| = \lambda k.$$

Thus, $|\text{supp}(u^n) \cup \text{supp}(v^n)| = (2-\lambda)k$.

A test separates u^n and v^n if it is positive for exactly one of the two vectors. This can happen in two ways. First, the test may be positive under u^n and negative under v^n . This means that the test avoids all infected items in v^n , but hits at least one infected item in $u^n \setminus v^n$. The probability of this event is

$$(1-p)^k - (1-p)^{(2-\lambda)k}.$$

Here, $(1-p)^k$ is the probability of avoiding all infected items in v^n , and $(1-p)^{(2-\lambda)k}$ is the probability of avoiding all infected items in $\text{supp}(u^n) \cup \text{supp}(v^n)$. Subtracting removes the case where the test also misses all infected items in $u^n \setminus v^n$. By symmetry, the probability that the test is positive under v^n and negative under u^n is also $(1-p)^k - (1-p)^{(2-\lambda)k}$. These two separating events are disjoint. Therefore, the probability that one test separates u^n and v^n is $2(1-p)^k - 2(1-p)^{(2-\lambda)k}$. Consequently, the probability that one test gives the same outcome under u^n and v^n is

$$1 - 2(1-p)^k + 2(1-p)^{(2-\lambda)k}.$$

Now choose $p = \ln(2)/k$. When $\lambda \leq \lambda_n^*$ we can say (under asymptotic approximations) that $(1-p)^k \approx \frac{1}{2}$ and $(1-p)^{(2-\lambda)k} \approx 2^{-(2-\lambda)}$. Hence the one-test consistency probability satisfies

$$1 - 2(1-p)^k + 2(1-p)^{(2-\lambda)k} = 2^{-(1-\lambda)}.$$

Since tests are independent under the Bernoulli design, the probability that all T tests give the same outcome under u^n and v^n is at most $(2^{-(1-\lambda)})^T$. Therefore,

$$\mathbb{P}(Y(v^n) = Y(u^n) | \mathcal{O}_\lambda(u^n, v^n)) \leq 2^{-T(1-\lambda)[1-o_\epsilon(1)-o_n(1)]}.$$

A full technical proof is deferred to the longer version.

Combining Theorems 3 and 4 gives the following small-overlap contribution.

$$\begin{aligned} & \sum_{u^n \in T_\epsilon^{(n)}} p_T(u^n) \sum_{\lambda \leq \lambda_n^*} N(\lambda | u^n) \mathbb{P}(Y(v^n) = Y(u^n) | \mathcal{O}_\lambda) \\ & \stackrel{(a)}{\leq} \sum_{u^n \in T_\epsilon^{(n)}} p_T(u^n) \sum_{\lambda \leq \lambda_n^*} 2^{nq \log n [\beta(1-\theta)(1-\lambda)]} 2^{-T(1-\lambda)} \\ & \stackrel{(b)}{\leq} \sum_{\lambda \leq \lambda_n^*} 2^{-(1-\lambda)[T-\beta(1-\theta)nq \log n][1-\epsilon-o_n(1)]} \\ & \leq (nq+1)2^{-(1-\lambda_n^*)[T-\beta(1-\theta)nq \log n][1-o_\epsilon-o_n(1)]} \rightarrow 0, \end{aligned} \quad (14)$$

as $n \rightarrow \infty$, provided

$$T > \beta(1-\theta)nq \log n. \quad (15)$$

Here, (a) uses the Lemmas 3 and 4 and (b) is due to the fact

that $\sum_{u^n \in T_\epsilon^{(n)}} p_T(u^n) \leq 1$.

It remains to control the large-overlap term. For this part, we directly use the argument from [9] for the near-constant-column design.

Theorem 5 (Large-overlap error bound): Let $\lambda_n^* = 1 - 1/\log n$ and choose $\nu = \ln 2$. For every $u^n \in T_\epsilon^{(n)}$,

$$\mathbb{P}(\mathcal{C}_{>\lambda_n^*}(u^n) \mid U^n = u^n) \leq n^{\theta - \frac{T \ln 2}{nq \log n} + o(1)}.$$

Consequently, if

$$T > (1 + \delta) \frac{\theta}{\ln 2} nq \log n$$

for some fixed $\delta > 0$, then

$$\sum_{u^n \in T_\epsilon^{(n)}} p_T(u^n) \mathbb{P}(\mathcal{C}_{>\lambda_n^*}(u^n) \mid U^n = u^n) \rightarrow 0.$$

Above the threshold, each infected item is isolated in many tests with high probability. Hence, any competitor that removes such an item must add enough infected items to keep those tests positive. A large-overlap competitor changes too few infected coordinates to do this, so any consistent competitor must have overlap at most λ_n^* . Thus, the large-overlap error is ruled out by the argument of [1]. We revisit this in a longer version of the paper.

Combining the small-overlap and large-overlap bounds shows that

$$\tau^* < \max \left\{ \beta(1 - \theta), \frac{\theta}{\ln 2} \right\}.$$

This proves the achievability part of Theorem 3.

V. CONVERSE

We now prove the converse part of Theorem 3. The lower bound has two components. The first comes from the counting argument revisited in [19], and the second follows [1].

The lower bound corresponding to the first argument in the maximum in (16) follows from the source-coding bound of [23]. The bound gives

$$\tau^* \geq \lim_{n \rightarrow \infty} \frac{H(U^n)}{nq \log n}, \quad (16)$$

where $H(U^n)$ is the entropy of the infection vector. This entropy can be computed as [19], yielding

$$\tau^* \geq \beta(1 - \theta). \quad (17)$$

It remains to obtain the second term. Fix an arbitrary non-adaptive test design $\mathbf{X} \in \{0, 1\}^{T \times n}$ and a true infection vector u^n . An infected item i is called *disguised* if every test containing i also contains some other infected item. For $b \in \{0, 1\}$, define the set of disguised items with status b by

$$V_b^+(\mathbf{X}, u^n) := \left\{ i : \begin{array}{l} u_i = b, \forall t \in [T] \text{ with } X_{ti} = 1, \\ \exists j \neq i \text{ such that } X_{tj} u_j = 1 \end{array} \right\}. \quad (18)$$

Thus, every test containing an item in $V_b^+(\mathbf{X}, u^n)$ is already positive because of some infected item other than that item itself. Thus, changing the status of a disguised item does not change any test outcome by itself, because all tests containing it are already masked by other infected items.

The key observation is that each pair

$$(i, j) \in V_1^+(\mathbf{X}, u^n) \times V_0^+(\mathbf{X}, u^n)$$

generates another infection vector with the same number of infected items and the same test outcome. Define v^n by removing the infected disguised item i and adding the healthy disguised item j such that $v_i = 0, v_j = 1, v_\ell = u_\ell$ for all $\ell \notin \{i, j\}$. Then v^n has the same Hamming weight as u^n . Moreover, every test that contains i still remains positive because i was disguised, and every test that contains j was already positive because j was disguised. All other tests are unchanged. Therefore $Y(v^n) = Y(u^n)$. Consequently, if both $V_1^+(\mathbf{X}, u^n)$ and $V_0^+(\mathbf{X}, u^n)$ are large, then the decoder has many competing vectors that are indistinguishable from the true vector.

The following statement is the form of the Coja-Oghlan et al. [1] argument that we need.

Theorem 6 (Disguised individuals): For any $\epsilon > 0$, there exists $\theta_0 = \theta_0(\epsilon) < 1$ such that for every $\theta \in (\theta_0, 1)$, there exists $n_0 = n_0(\theta, \epsilon)$ such that for all $n > n_0$ and all non-adaptive test designs $\mathbf{X} \in \{0, 1\}^{T \times n}$ with

$$\tau < \frac{\theta}{\ln 2},$$

we have

$$\mathbb{P}(|V_0^+(\mathbf{X}, U^n)| |V_1^+(\mathbf{X}, U^n)| \geq \log n) > 1 - \epsilon.$$

Proof Sketch: The proof follows the disguised-individual argument of [1]. The only additional point is that the infection vector is sampled from a Markov source rather than uniformly from all k -subsets. For typical Markov sequences, however, the Hamming weight is $nq + o(nq)$ and the infected coordinates are sufficiently spread over the tests for the same masking calculation to apply. We consider this more precisely in a longer version of the paper.

The result above gives the disguised-individual lower bound for θ sufficiently close to 1. The reduction argument in [1] then extends this lower bound to all

$$\theta > \frac{\ln 2}{1 + \ln 2}.$$

Therefore, throughout this regime, there is a non-vanishing probability that at least one other typical vector produces exactly the same test outcome as the true vector under any sequence of test matrices and decoders.

Combining the counting lower bound and the disguised-individual lower bound gives

$$\tau^* \geq \max \left\{ \beta(1 - \theta), \frac{\theta}{\ln 2} \right\}.$$

Together with the achievability result, this proves Theorem 3.

VI. CONCLUSION

We characterized the non-adaptive testing capacity under a sparse Markovian infection prior. Although the main result is stated for a two-state Markov model, the proof techniques are substantially more general. In particular, for any fixed k -th order Markov source, the same argument can be used.

More broadly, the analysis relies on two ingredients: typical-set counting for the source prior, and overlap-based indistinguishability bounds for the test design. These ingredients are not specific to the Markov model and provide a natural route for studying group testing under broader correlated priors. Thus, the work shows how structural properties of the infection prior can be translated into sharp group testing thresholds.

REFERENCES

- [1] A. Coja-Oghlan, O. Gebhard, M. Hahn-Klimroth, and P. Loick, “Optimal group testing,” in *Conference on Learning Theory*. PMLR, 2020, pp. 1374–1388.
- [2] D. J. Balding, W. J. Bruno, D. C. Torney, and E. Knill, “A comparative survey of non-adaptive pooling designs,” in *Genetic mapping and DNA sequencing*. Springer, 1996, pp. 133–154.
- [3] R. Dorfman, “The detection of defective members of large populations,” *The Annals of mathematical statistics*, vol. 14, no. 4, pp. 436–440, 1943.
- [4] M. Aldridge, O. Johnson, J. Scarlett *et al.*, “Group testing: an information theory perspective,” *Foundations and Trends® in Communications and Information Theory*, vol. 15, no. 3–4, pp. 196–392, 2019.
- [5] T. Li, C. L. Chan, W. Huang, T. Kaced, and S. Jaggi, “Group testing with prior statistics,” in *2014 IEEE International Symposium on Information Theory*. IEEE, 2014, pp. 2346–2350.
- [6] J. Scarlett and V. Cevher, “Phase transitions in the pooled data problem,” *Advances in Neural Information Processing Systems*, vol. 30, 2017.
- [7] M. Aldridge, “The capacity of bernoulli nonadaptive group testing,” *IEEE Transactions on Information Theory*, vol. 63, no. 11, pp. 7142–7148, 2017.
- [8] T. Berger and V. I. Levenshtein, “Asymptotic efficiency of two-stage disjunctive testing,” *IEEE Transactions on Information Theory*, vol. 48, no. 7, pp. 1741–1749, 2002.
- [9] A. Coja-Oghlan, O. Gebhard, M. Hahn-Klimroth, and P. Loick, “Information-theoretic and algorithmic thresholds for group testing,” *IEEE Transactions on Information Theory*, vol. 66, no. 12, pp. 7911–7928, 2020.
- [10] C. J. Colbourn, “Group testing for consecutive positives,” *Annals of Combinatorics*, vol. 3, no. 1, pp. 37–41, 1999.
- [11] T. V. Bui, M. Cheraghchi, and T. D. Nguyen, “Improved algorithms for non-adaptive group testing with consecutive positives,” in *2021 IEEE International Symposium on Information Theory (ISIT)*. IEEE, 2021, pp. 1961–1966.
- [12] Y.-H. Li, R. Gabrys, J. Sima, I. Shomorony, and O. Milenkovic, “Finding a burst of positives via nonadaptive semiquantitative group testing,” in *2023 IEEE International Symposium on Information Theory (ISIT)*. IEEE, 2023, pp. 1848–1853.
- [13] A. Janesick, R. Shelansky, A. D. Gottscho, F. Wagner, S. R. Williams, M. Rouault, G. Beliakoff, C. A. Morrison, M. F. Oliveira, J. T. Sichertman *et al.*, “High resolution mapping of the tumor microenvironment using integrated single-cell, spatial and in situ analysis,” *Nature Communications*, vol. 14, no. 1, p. 8353, 2023.
- [14] P. Nikolopoulos, S. R. Srinivasavaradhan, T. Guo, C. Fragouli, and S. Diggavi, “Group testing for connected communities,” in *International Conference on Artificial Intelligence and Statistics*. PMLR, 2021, pp. 2341–2349.
- [15] P. Nikolopoulos, S. R. Srinivasavaradhan, T. Guo, C. Fragouli, and S. N. Diggavi, “Community-aware group testing,” *IEEE Transactions on Information Theory*, vol. 69, no. 7, pp. 4361–4383, 2023.
- [16] S. Ahn, W.-N. Chen, and A. Özgür, “Noisy adaptive group testing for community-oriented models,” in *2023 IEEE International Symposium on Information Theory (ISIT)*. IEEE, 2023, pp. 1621–1626.
- [17] —, “Adaptive group testing on networks with community structure,” in *2021 IEEE International Symposium on Information Theory (ISIT)*. IEEE, 2021, pp. 1242–1247.
- [18] B. Arasli and S. Ulukus, “Group testing with a graph infection spread model,” *Information*, vol. 14, no. 1, p. 48, 2023.
- [19] A. N. Ravi and I. Shomorony, “Fundamental limits of non-adaptive group testing with markovian correlation,” *arXiv preprint*, vol. <https://arxiv.org/abs/2501.12588>, 2025, preprint.
- [20] A. C. Portnoy and A. Cohen, “Non-adaptive multi-stage algorithm for group testing with prior statistics,” in *2024 60th Annual Allerton Conference on Communication, Control, and Computing*. IEEE, 2024, pp. 1–8.
- [21] I. Lau, J. Scarlett, and Y. Sun, “Model-based and graph-based priors for group testing,” *IEEE Transactions on Signal Processing*, vol. 70, pp. 6035–6050, 2022.
- [22] H. Nikpey, S. Sarkar, and S. S. Bidokhti, “Group testing with general correlation using hypergraphs,” in *2024 IEEE International Symposium on Information Theory (ISIT)*. IEEE, 2024, pp. 3225–3230.
- [23] J. Wolf, “Born again group testing: Multiaccess communications,” *IEEE Transactions on Information Theory*, vol. 31, no. 2, pp. 185–191, 1985.